

Task Overview :-Set up a data pipeline that extracts data from a MongoDB database and syncs it into Apache Iceberg using OLake's Hive integration, then query it using Apache Spark.

Steps:

1. Spin up a local MongoDB database (via Docker or natively).
2. Create a sample table (e.g., orders) and load it with ~10-20 records.
3. Use OLake's discover and sync commands to:
 - Discover schema from MongoDB
 - Sync data into Iceberg format via Hive Metastore
4. Query synced Iceberg tables using Spark SQL (either via CLI or Spark UI).

Submission:

- Loom video walkthrough (2–5 min)
- Submit a Markdown report (README.md) that includes:
 - Setup steps
 - Screenshots of Spark UI or SQL results
 - Any challenges faced or improvements you'd suggest

Resources:

- OLake GitHub– <https://github.com/datazip-inc/olake>
- OLake Docs - <https://olake.io/docs/> (you can find almost everything required here)
- Iceberg Docs - <https://iceberg.apache.org/docs/latest/>
- Spark Docs -<https://spark.apache.org/docs/latest/>